

benchmark_analysis_concise

May 3, 2025

1 LucidBench Benchmark Analysis

This notebook provides comprehensive analysis of filesystem benchmark results across different storage devices and filesystems.

1.1 Table of Contents

1. Setup and Data Loading
2. Overview and Summary Statistics
3. Performance Analysis by Storage Type
4. Filesystem Comparison
5. I/O Pattern Analysis
6. Resource Utilization
7. Statistical Analysis
8. **Comparative Analysis**
9. Recommendations and Conclusions

1.2 1. Setup and Data Loading

```
[2]: import os
import json
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
from scipy import stats
from pathlib import Path

# Set plot style
plt.style.use('seaborn-v0_8')
sns.set_palette('husl')
# Define consistent colors for filesystems
fs_colors = {
    'ext4': '#1f77b4', # blue
    'ext3': '#ff7f0e', # orange
    'ext2': '#2ca02c', # green
    'xfs': '#d62728', # red
```

```

    'btrfs': '#9467bd', # purple
    'ntfs': '#8c564b', # brown
    'vfat': '#e377c2' # pink
}

# Configure display options
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 100)
pd.set_option('display.width', 1000)

```

```

[3]: def load_benchmark_data(run_dir):
    """Load all benchmark and monitoring data from a run directory."""
    data = []

    for test_dir in Path(run_dir).glob('*'):
        if not test_dir.is_dir():
            continue

        # Parse test directory name
        parts = test_dir.name.split('_')
        if len(parts) < 4:
            continue

        storage_type = parts[0]
        device = parts[1]
        filesystem = parts[2]
        test_type = '_'.join(parts[3:])

        # Load test data
        test_file = test_dir / 'test.json'
        if test_file.exists():
            with open(test_file) as f:
                content = f.read()
                start_idx = content.find('{')
                if start_idx >= 0:
                    test_data = json.loads(content[start_idx:])
                else:
                    test_data = json.loads(content)

        # Load monitoring data
        monitor_file = test_dir / 'monitoring.json'
        if monitor_file.exists():
            with open(monitor_file) as f:
                monitor_data = json.load(f)

        # Combine data

```

```

        if 'test_data' in locals() and 'monitor_data' in_
↪locals():
            data.append({
                'storage_type': storage_type,
                'device': device,
                'filesystem': filesystem,
                'test_type': test_type,
                'test_data': test_data,
                'monitor_data': monitor_data
            })

    return pd.DataFrame(data)

# Load data from the most recent run
results_dir = Path('../results')
latest_run = max(results_dir.glob('run_*'), key=os.path.getctime)
df = load_benchmark_data(latest_run)
print(f"Loaded data from {latest_run.name}")

```

Loaded data from run_20250427_234135

1.3 2. Overview and Summary Statistics

```

[4]: def extract_performance_metrics(row):
    """Extract key performance metrics from test data."""
    test_data = row['test_data']
    job = test_data['jobs'][0]
    test_type = row['test_type']

    # Determine if this is a read or write test
    is_read_test = 'read' in test_type

    metrics = {
        'iops': job['read']['iops'] if is_read_test else job['write']['iops'],
        'bandwidth': job['read']['bw'] if is_read_test else job['write']['bw'],
        'latency': job['read']['lat_ns']['mean'] if is_read_test else_
↪job['write']['lat_ns']['mean'],
        'runtime': job.get('runtime', None)
    }

    return pd.Series(metrics)

# Extract performance metrics
performance_df = df.apply(extract_performance_metrics, axis=1)
df = pd.concat([df, performance_df], axis=1)

# Display summary statistics

```

```
print("\nSummary Statistics by Storage Type and Filesystem:")
summary = df.groupby(['storage_type', 'filesystem']).agg({
    'iops': ['mean', 'std', 'min', 'max'],
    'bandwidth': ['mean', 'std', 'min', 'max'],
    'latency': ['mean', 'std', 'min', 'max']
}).round(2)
display(summary)
```

Summary Statistics by Storage Type and Filesystem:

		iops				bandwidth		
		latency						
		mean	std	min	max	mean		
std	min	max	mean	std	min	max	min	
max								
storage_type	filesystem							
HDD	btrfs	243.20	31.97	224.37	291.05	117599.75		
↳134603.50	897.0	234764.0	1.330585e+08	1.547450e+07	1.099335e+08	1.		
↳426070e+08								
	ext2	175.52	79.58	100.37	259.58	84171.50		
↳109908.52	460.0	232410.0	2.143106e+08	9.661486e+07	1.232628e+08	3.		
↳165003e+08								
	ext3	175.02	78.72	100.09	257.37	84067.00		
↳109836.12	463.0	232281.0	2.144025e+08	9.619820e+07	1.243215e+08	3.		
↳171924e+08								
	ext4	230.31	34.30	208.26	281.41	110971.75		
↳127023.98	833.0	223374.0	1.409504e+08	1.838856e+07	1.137028e+08	1.		
↳536348e+08								
	ntfs	9388.24	17002.00	133.57	34873.49	686754.75		
↳793607.88	534.0	1726761.0	7.407539e+07	1.112951e+08	9.158244e+05	2.		
↳395449e+08								
	vfat	147.35	117.02	0.00	274.80	30933.00		
↳60613.76	0.0	121851.0	1.369791e+08	1.110981e+08	0.000000e+00	2.		
↳678751e+08								
	xfs	241.13	33.68	214.26	290.43	118222.75		
↳135346.89	857.0	235582.0	1.344233e+08	1.687346e+07	1.101695e+08	1.		
↳493296e+08								
NVMe	btrfs	24226.72	39104.99	1633.17	82617.08	1190281.00		
↳1242254.81	40008.0	2718278.0	8.778315e+06	8.715936e+06	3.864838e+05	1.		
↳955125e+07								
	ext2	41196.32	58379.08	1352.26	125547.89	1184340.00		
↳1143129.40	140956.0	2709498.0	9.170515e+06	1.097653e+07	2.540123e+05	2.		
↳350814e+07								
	ext3	39995.97	57055.94	1240.08	122554.46	1103520.25		
↳1056072.32	134917.0	2519101.0	9.943538e+06	1.197304e+07	2.602802e+05	2.		
↳564301e+07								

	ext4	46855.52	58478.42	1913.12	125367.77	1332587.50	1.
↳1155650.63	230253.0	2639587.0	7.463895e+06	8.344051e+06	2.544177e+05	1.	
↳670154e+07							
	ntfs	9796.01	17711.95	127.51	36343.27	731068.50	2.
↳860525.91	510.0	1878326.0	7.631343e+07	1.173010e+08	8.791321e+05	2.	
↳509305e+08							
	vfat	35424.96	43129.55	0.00	88323.45	261890.75	6.
↳206723.22	0.0	482651.0	1.713611e+07	3.362955e+07	0.000000e+00	6.	
↳757906e+07							
	xfs	38455.01	57503.74	1840.07	123014.55	1300392.00	1.
↳1216288.80	105236.0	2720041.0	7.700020e+06	8.346075e+06	2.592703e+05	1.	
↳736146e+07							
SSD	btrfs	7820.49	9191.93	472.71	19544.02	282023.75	6.
↳256573.54	43018.0	522850.0	3.370955e+07	3.632652e+07	1.635174e+06	6.	
↳764454e+07							
	ext2	8780.89	9840.12	378.73	19586.37	259106.75	8.
↳225977.00	58635.0	511625.0	3.798088e+07	4.246213e+07	1.631772e+06	8.	
↳414086e+07							
	ext3	8755.49	9816.28	368.01	19533.83	256287.25	8.
↳224075.12	58481.0	511687.0	3.859024e+07	4.335040e+07	1.636321e+06	8.	
↳658831e+07							
	ext4	8391.63	9188.55	399.80	16600.85	256014.75	7.
↳222208.72	64373.0	483883.0	3.788818e+07	4.179290e+07	1.925732e+06	7.	
↳996131e+07							
	ntfs	9761.84	17687.43	130.27	36272.87	713318.75	2.
↳846692.76	521.0	1853426.0	7.554200e+07	1.144046e+08	8.806148e+05	2.	
↳456175e+08							
	vfat	9016.33	10386.80	0.00	19611.28	87070.00	1.
↳85648.63	0.0	204820.0	4.076104e+07	7.912891e+07	0.000000e+00	1.	
↳594474e+08							
	xfs	9095.30	10038.48	483.08	19507.66	289783.25	6.
↳253221.60	63519.0	522915.0	3.310901e+07	3.615384e+07	1.638377e+06	6.	
↳620603e+07							

1.4 3. Performance Analysis by Storage Type


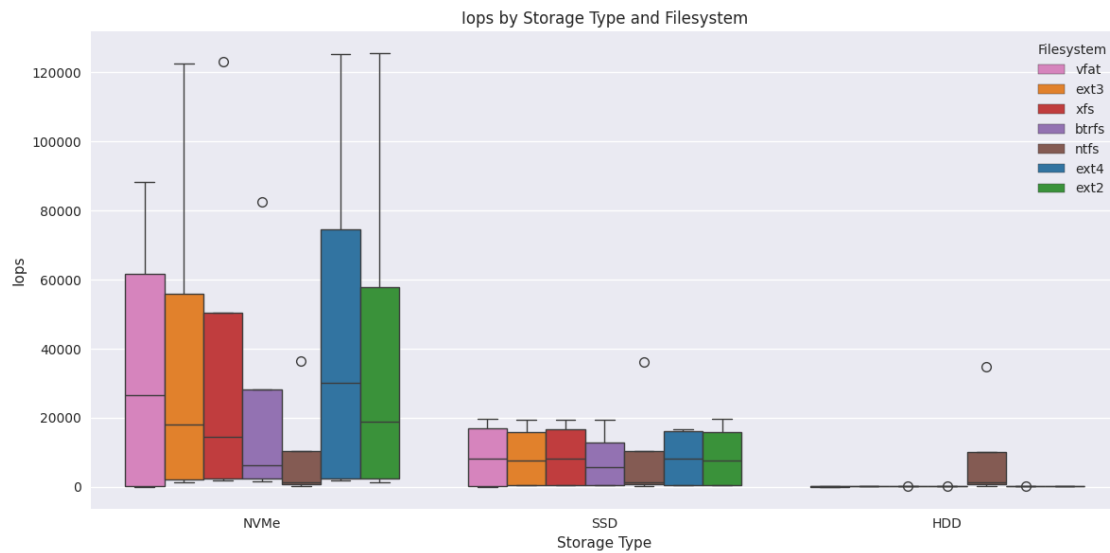
```
[5]: def plot_storage_performance(df, metric):
    """Plot performance metrics by storage type."""
    plt.figure(figsize=(12, 6))

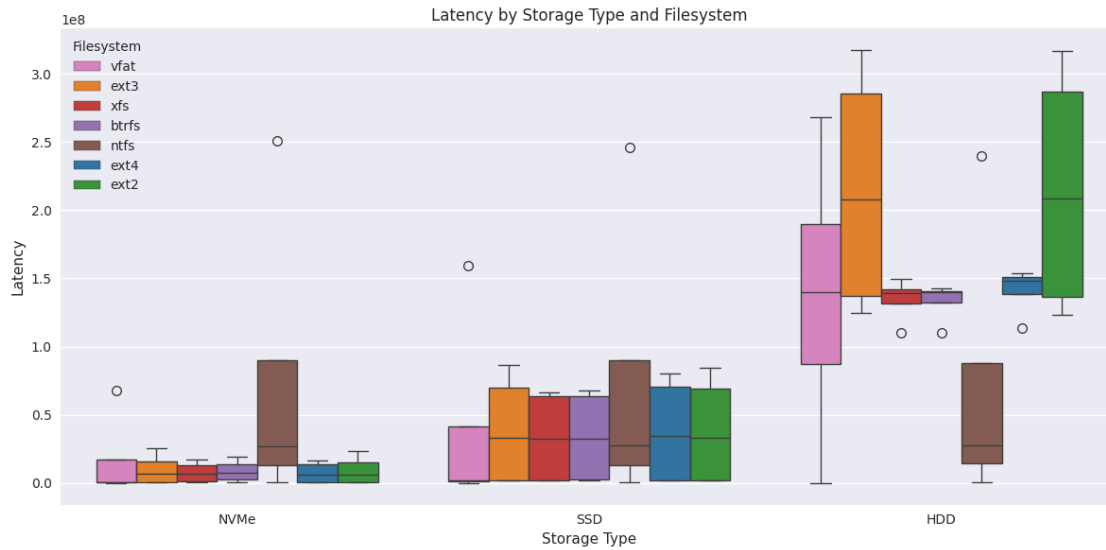
    # Create box plot
    sns.boxplot(data=df, x='storage_type', y=metric, hue='filesystem',
↳palette=fs_colors)

    plt.title(f'{metric.title()} by Storage Type and Filesystem')
    plt.xlabel('Storage Type')
    plt.ylabel(metric.title())
```

```
plt.xticks(rotation=0)
plt.legend(title='Filesystem')
plt.tight_layout()
plt.show()

# Plot performance metrics
for metric in ['iops', 'bandwidth', 'latency']:
    plot_storage_performance(df, metric)
```





1.5 4. Filesystem Comparison

```
[6]: def analyze_filesystem_performance(df):
    """Analyze filesystem performance across different test types."""
    # Create pivot table for each metric
    metrics = ['iops', 'bandwidth', 'latency']

    for metric in metrics:
        pivot = pd.pivot_table(
            df,
            values=metric,
            index=['storage_type', 'test_type'],
            columns='filesystem',
            aggfunc='mean'
        )

        print(f"\n{metric.upper()} Comparison:")
        display(pivot)

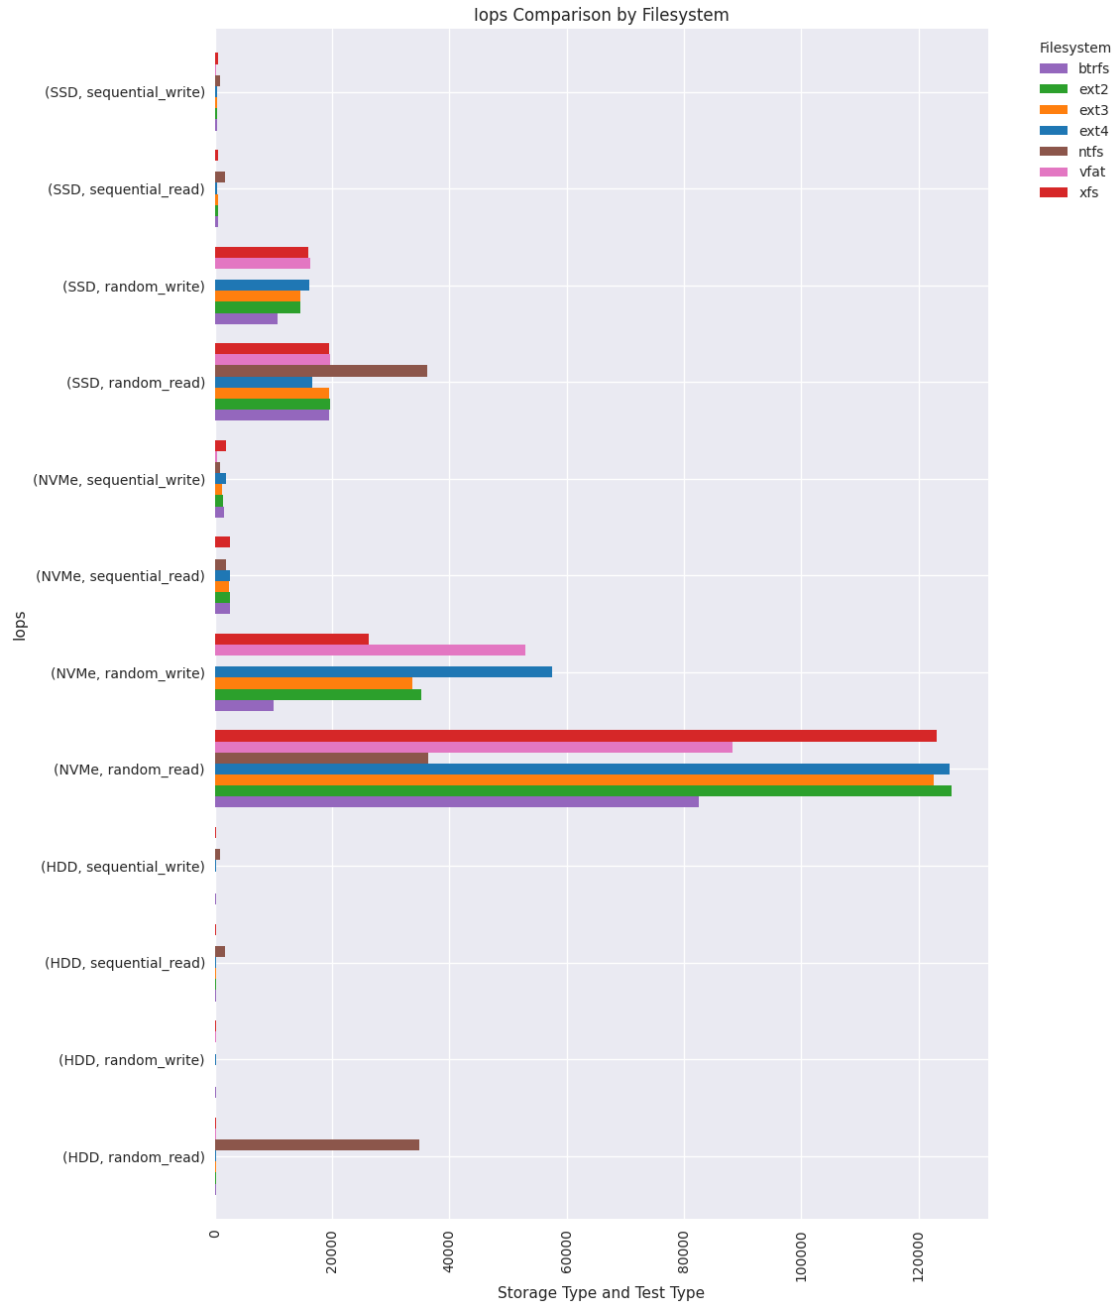
        # Plot comparison
        fig, ax = plt.subplots(figsize=(12, 15))
        # Use a distinct color palette with high contrast colors
        pivot.plot(kind='barh', width=0.8, ax=ax, color=fs_colors) # Increased
        # bar width and custom colors
        plt.title(f'{metric.title()} Comparison by Filesystem')
        plt.xlabel('Storage Type and Test Type')
        plt.ylabel(metric.title())
        plt.xticks(rotation=90)
```

```
plt.legend(title='Filesystem', bbox_to_anchor=(1.05, 1), loc='upper_
↳left')
# Ensure plot takes up full figure size while leaving room for legend
plt.subplots_adjust(left=0.2, right=0.85, top=0.9, bottom=0.1)
plt.show()
```

```
analyze_filesystem_performance(df)
```

IOPS Comparison:

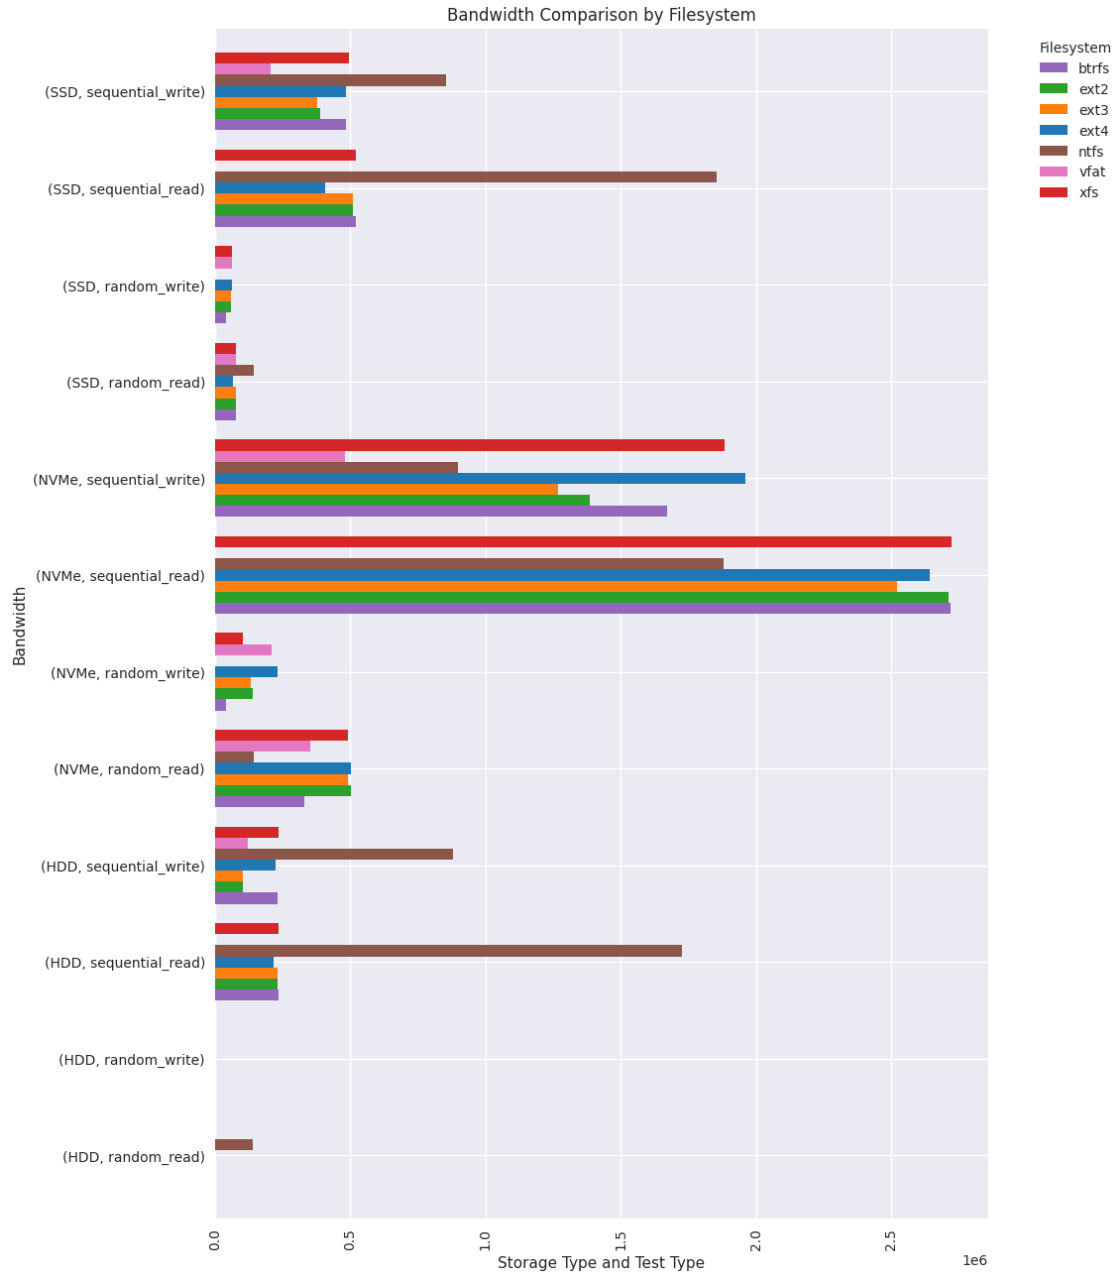
filesystem		btrfs	ext2	ext3	
↳ ext4	ntfs	vfat	xfs		
storage_type	test_type				
HDD	random_read	291.049988	259.577509	257.369935	↳
↳281.408540	34873.486763	274.804210	290.429866		
	random_write	224.369636	115.179494	115.785790	↳
↳208.261990	133.572202	195.559394	214.259616		
	sequential_read	229.262286	226.962930	226.837238	↳
↳213.433380	1686.290655	0.000000	229.776731		
	sequential_write	228.100462	100.370016	100.092860	↳
↳218.139213	859.601259	119.024787	230.060661		
NVMe	random_read	82617.081626	125547.892720	122554.464703	↳
↳125367.766619	36343.269097	88323.450135	123014.547161		
	random_write	10002.060361	35239.145046	33729.284611	↳
↳57563.460694	127.512081	52904.944501	26309.112806		
	sequential_read	2654.569021	2645.994832	2460.060060	↳
↳2577.721838	1834.303627	0.000000	2656.290532		
	sequential_write	1633.173844	1352.261472	1240.084771	↳
↳1913.124708	878.969957	471.454880	1840.071878		
SSD	random_read	19544.024454	19586.371787	19533.830104	↳
↳16600.848585	36272.865643	19611.281514	19507.664831		
	random_write	10754.625641	14658.838003	14620.412716	↳
↳16093.314507	130.271774	16253.968254	15879.815847		
	sequential_read	510.595861	499.634057	499.695010	↳
↳399.804783	1809.986743	0.000000	510.659519		
	sequential_write	472.706290	378.733241	368.014376	↳
↳472.542686	834.215886	200.068383	483.075834		



BANDWIDTH Comparison:

filesystem	btrfs	ext2	ext3	ext4	
ntfs					
vfat					
xfs					
storage_type	test_type				
HDD	random_read	1164.0	1038.0	1029.0	1125.0
		139493.0	1099.0	1161.0	


	random_write	897.0	460.0	463.0	833.0	┐
↪534.0	782.0 857.0					
	sequential_read	234764.0	232410.0	232281.0	218555.0	┐
↪1726761.0	0.0 235291.0					
	sequential_write	233574.0	102778.0	102495.0	223374.0	┐
↪880231.0	121851.0 235582.0					
NVMe	random_read	330468.0	502191.0	490217.0	501471.0	┐
↪145373.0	353293.0 492058.0					
	random_write	40008.0	140956.0	134917.0	230253.0	┐
↪510.0	211619.0 105236.0					
	sequential_read	2718278.0	2709498.0	2519101.0	2639587.0	┐
↪1878326.0	0.0 2720041.0					
	sequential_write	1672370.0	1384715.0	1269846.0	1959039.0	┐
↪900065.0	482651.0 1884233.0					
SSD	random_read	78176.0	78345.0	78135.0	66403.0	┐
↪145091.0	78445.0 78030.0					
	random_write	43018.0	58635.0	58481.0	64373.0	┐
↪521.0	65015.0 63519.0					
	sequential_read	522850.0	511625.0	511687.0	409400.0	┐
↪1853426.0	0.0 522915.0					
	sequential_write	484051.0	387822.0	376846.0	483883.0	┐
↪854237.0	204820.0 494669.0					



LATENCY Comparison:

filesystem	ntfs	vfat	btrfs	xfs	ext2	ext3	
ext4							
storage_type	test_type						
HDD	random_read	1.099335e+08	1.232628e+08	1.243215e+08	1.		
		137028e+08	9.158244e+05	1.164325e+08	1.101695e+08		

	random_write	1.426070e+08	2.765953e+08	2.751344e+08	1.
↪536348e+08		2.395449e+08	1.636089e+08	1.493296e+08	
	sequential_read	1.394754e+08	1.408839e+08	1.409617e+08	1.
↪498208e+08		1.878933e+07	0.000000e+00	1.391576e+08	
	sequential_write	1.402183e+08	3.165003e+08	3.171924e+08	1.
↪466432e+08		3.705144e+07	2.678751e+08	1.390365e+08	
NVMe	random_read	3.864838e+05	2.540123e+05	2.602802e+05	2.
↪544177e+05		8.791321e+05	3.616420e+05	2.592703e+05	
	random_write	3.195223e+06	9.062730e+05	9.469127e+05	5.
↪548055e+05		2.509305e+08	6.037502e+05	1.213946e+06	
	sequential_read	1.198030e+07	1.201364e+07	1.292395e+07	1.
↪234481e+07		1.727433e+07	0.000000e+00	1.196540e+07	
	sequential_write	1.955125e+07	2.350814e+07	2.564301e+07	1.
↪670154e+07		3.616977e+07	6.757906e+07	1.736146e+07	
SSD	random_read	1.635174e+06	1.631772e+06	1.636321e+06	1.
↪925733e+06		8.806149e+05	1.630163e+06	1.638377e+06	
	random_write	2.973479e+06	2.180709e+06	2.186319e+06	1.
↪986038e+06		2.456175e+08	1.966600e+06	2.012452e+06	
	sequential_read	6.258500e+07	6.397019e+07	6.395002e+07	7.
↪996131e+07		1.750670e+07	0.000000e+00	6.257917e+07	
	sequential_write	6.764454e+07	8.414086e+07	8.658831e+07	6.
↪767964e+07		3.816323e+07	1.594474e+08	6.620603e+07	



1.6 5. I/O Pattern Analysis

```
[7]: def analyze_io_patterns(df):
    """Analyze I/O patterns across different test types."""
    # Group by test type and calculate statistics
    io_stats = df.groupby('test_type').agg({
        'iops': ['mean', 'std', 'min', 'max'],
        'bandwidth': ['mean', 'std', 'min', 'max'],
```

```

        'latency': ['mean', 'std', 'min', 'max']
    }).round(2)

print("\nI/O Pattern Statistics:")
display(io_stats)

# Plot I/O patterns
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# IOPS by test type
sns.boxplot(data=df, x='test_type', y='iops', ax=axes[0,0])
axes[0,0].set_title('IOPS by Test Type')
axes[0,0].set_xticks(range(len(df['test_type'].unique())))
axes[0,0].set_xticklabels(df['test_type'].unique(), rotation=45)

# Bandwidth by test type
sns.boxplot(data=df, x='test_type', y='bandwidth', ax=axes[0,1])
axes[0,1].set_title('Bandwidth by Test Type')
axes[0,1].set_xticks(range(len(df['test_type'].unique())))
axes[0,1].set_xticklabels(df['test_type'].unique(), rotation=45)

# Latency by test type
sns.boxplot(data=df, x='test_type', y='latency', ax=axes[1,0])
axes[1,0].set_title('Latency by Test Type')
axes[1,0].set_xticks(range(len(df['test_type'].unique())))
axes[1,0].set_xticklabels(df['test_type'].unique(), rotation=45)

# Create a more meaningful visualization showing relationship between
↪ metrics
# Use a heatmap to show correlation between key performance metrics
correlation = df[['iops', 'bandwidth', 'latency']].corr()
sns.heatmap(correlation, annot=True, cmap='coolwarm', vmin=-1, vmax=1,
↪ ax=axes[1,1])
axes[1,1].set_title('Correlation Between Performance Metrics')

plt.tight_layout()
plt.show()

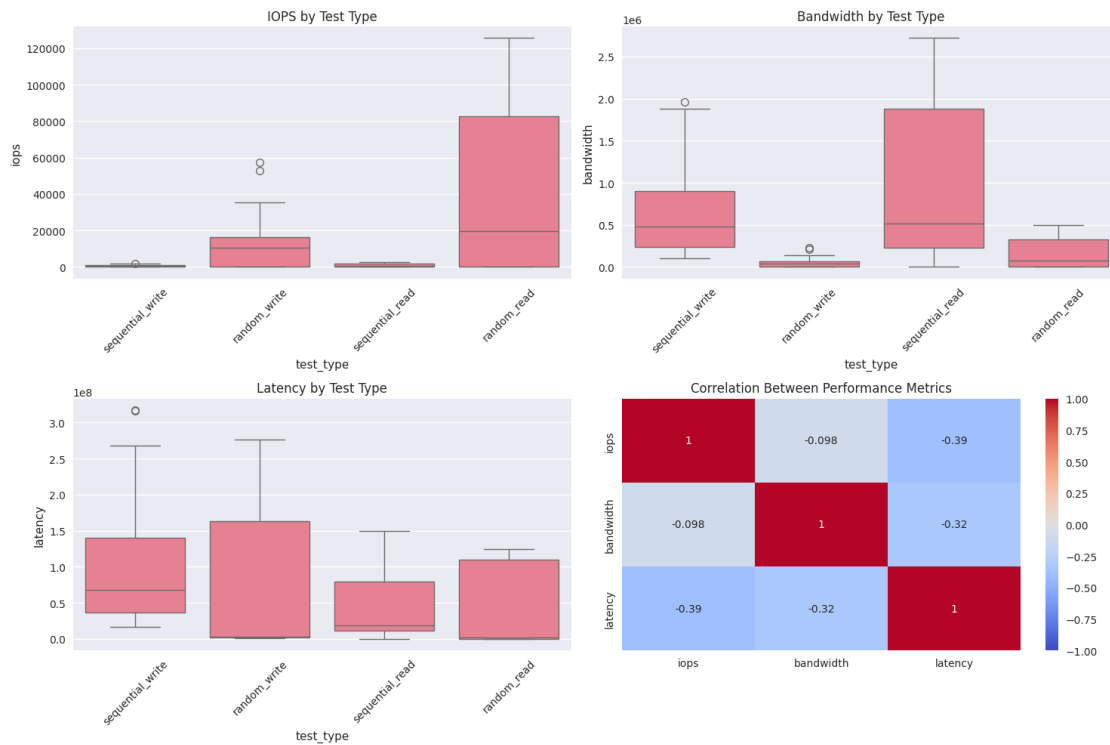
analyze_io_patterns(df)

```

I/O Pattern Statistics:

	iops				bandwidth			
	mean		std		min		max	
↪ test_type	min	max	mean	std	min	max	mean	std

random_read	42426.36	47169.54	257.37	125547.89	169705.00	188678.13
↪	1029.0	502191.0	3.392246e+07	5.348442e+07	254012.28	1.243215e+08
random_write	14546.37	17550.69	115.18	57563.46	58185.10	70202.60
↪	460.0	230253.0	9.132045e+07	1.119830e+08	554805.46	2.765953e+08
sequential_read	1041.52	1047.36	0.00	2656.29	1066514.10	1072498.11
↪	0.0	2720041.0	5.514969e+07	5.533024e+07	0.00	1.498208e+08
sequential_write	685.42	584.35	100.09	1913.12	701863.43	598375.13
↪	102495.0	1959039.0	1.019477e+08	9.465585e+07	16701543.24	3.171924e+08



1.7 6. Resource Utilization

1.7.1 Analytics based on the monitoring data

This section analyzes CPU, memory, and disk I/O utilization patterns across different storage types and filesystems. Note that the NTFS results show significantly higher resource utilization due to its much longer elapsed time (7132 seconds) compared to other filesystems (20-70 seconds), which may skew the overall resource utilization metrics.

```
[8]: def plot_resource_trends(df, metric, title):
    """
    Plot resource usage trends over time for different storage types and
    ↪filesystems.

    Args:
```

```

        df: DataFrame containing the monitoring data
        metric: Metric to plot (e.g., 'cpu_percent', 'memory_percent')
        title: Title for the plot
        """
        # Create subplots for each storage type
        storage_types = df['storage_type'].unique()
        fig, axes = plt.subplots(len(storage_types), 1, figsize=(15, 3*len(storage_types)))

        for idx, storage_type in enumerate(storage_types):
            ax = axes[idx]

            # Filter data for current storage type
            storage_df = df[df['storage_type'] == storage_type]

            for filesystem, group in storage_df.groupby('filesystem'):
                if 'monitor_data' in group.iloc[0] and 'stats' in group.
                iloc[0]['monitor_data']:
                    stats = pd.DataFrame(group.iloc[0]['monitor_data']['stats'])
                    stats['timestamp'] = pd.to_datetime(stats['timestamp'])
                    ax.plot(range(len(stats)), stats[metric],
                            label=filesystem,
                            color=fs_colors[filesystem],
                            alpha=0.7)

            ax.set_title(f'{title} Usage Over Time - {storage_type}')
            ax.set_xlabel('Time (seconds)')
            ax.set_ylabel(title)
            ax.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
            ax.grid(True, alpha=0.3)

        plt.tight_layout()
        plt.show()

# Plot Disk I/O Throughput
def plot_io_throughput(df):
    """
    Plot disk I/O throughput over time for different storage types and
    filesystems.

    Args:
        df: DataFrame containing the monitoring data
    """
    fig, axes = plt.subplots(3, 1, figsize=(15, 9))
    storage_types = df['storage_type'].unique()

    for idx, storage_type in enumerate(storage_types):

```



```

ax = axes[idx]

    for filesystem, group in df[df['storage_type'] == storage_type].
↳groupby('filesystem'):
        if 'monitor_data' in group.iloc[0] and 'stats' in group.
↳iloc[0]['monitor_data']:
            stats = pd.DataFrame(group.iloc[0]['monitor_data']['stats'])
            stats['timestamp'] = pd.to_datetime(stats['timestamp'])

            # Calculate I/O throughput (MB/s)
            stats['io_throughput'] = (stats['disk_read_bytes'].diff() +
                                     stats['disk_write_bytes'].diff()) / 1024 / 1024

        ax.plot(range(len(stats)), stats['io_throughput'],
                label=filesystem,
                color=fs_colors[filesystem],
                alpha=0.7)

    ax.set_title(f'Disk I/O Throughput Over Time - {storage_type}')
    ax.set_xlabel('Time (seconds)')
    ax.set_ylabel('Throughput (MB/s)')
    ax.legend(bbox_to_anchor=(1.05, 1), loc='upper left')
    ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Plot CPU Usage
plot_resource_trends(df, 'cpu_percent', 'CPU')

# Plot Memory Usage
plot_resource_trends(df, 'memory_percent', 'Memory')

# Plot Disk I/O Throughput
plot_io_throughput(df)

# Resource usage summary statistics
resource_summary = pd.DataFrame()
for (storage_type, filesystem), group in df.groupby(['storage_type',
↳'filesystem']):
    if 'monitor_data' in group.iloc[0] and 'stats' in group.
↳iloc[0]['monitor_data']:
        stats = pd.DataFrame(group.iloc[0]['monitor_data']['stats'])
        # Convert timestamp strings to datetime objects
        stats['timestamp'] = pd.to_datetime(stats['timestamp'])
        summary = {

```


```

        'Storage Type': storage_type,
        'Filesystem': filesystem,
        'Avg CPU %': stats['cpu_percent'].mean(),
        'Max CPU %': stats['cpu_percent'].max(),
        'Avg Memory %': stats['memory_percent'].mean(),
        'Max Memory %': stats['memory_percent'].max(),
        'Total Read (GB)': stats['disk_read_bytes'].max() / 1024**3,
        'Total Write (GB)': stats['disk_write_bytes'].max() / 1024**3,
        'Elapsed Time (s)': (stats['timestamp'].max() - stats['timestamp'].
↪min()).total_seconds()
    }
    resource_summary = pd.concat([resource_summary, pd.
↪DataFrame([summary])], ignore_index=True)

display(resource_summary)

```





	Storage Type	Filesystem	Avg CPU %	Max CPU %	Avg Memory %	Max Memory %	
↩	Total Read (GB)		Total Write (GB)		Elapsed Time (s)		
0	HDD	btrfs	0.222857	2.3	1.408571	1.5	↩
↩	304.714440	3234.096130	68.070094				

1	HDD	ext2	0.127778	2.3	1.400000	1.4	␣
↪	291.600677	3161.113423		142.137219			
2	HDD	ext3	0.143478	2.6	1.400000	1.4	␣
↪	291.850449	3171.901868		182.176519			
3	HDD	ext4	0.306667	2.3	1.400000	1.4	␣
↪	300.175681	3222.764630		28.030823			
4	HDD	ntfs	1.918164	3.1	2.910205	3.5	␣
↪	287.261601	2553.800095		3117.673241			
5	HDD	vfat	0.332432	2.3	1.410811	1.5	␣
↪	285.962682	861.352174		72.074149			
6	HDD	xfs	0.272727	2.3	1.400000	1.4	␣
↪	300.442266	3224.521772		64.065516			
7	NVMe	btrfs	2.175000	13.7	1.500000	1.5	␣
↪	403.187676	18510.434410		22.021028			
8	NVMe	ext2	0.616667	3.1	1.500000	1.5	␣
↪	374.987468	18340.320305		34.032300			
9	NVMe	ext3	0.621053	2.2	1.500000	1.5	␣
↪	383.167312	18474.386311		36.033045			
10	NVMe	ext4	1.163636	7.5	1.500000	1.6	␣
↪	387.047649	18476.867687		20.020993			
11	NVMe	ntfs	1.032052	3.7	3.253597	3.5	␣
↪	366.742051	16343.229915		22656.318134			
12	NVMe	vfat	1.530000	2.6	1.400000	1.4	␣
↪	364.681422	10745.573076		18.020437			
13	NVMe	xfs	1.063636	7.1	1.518182	1.6	␣
↪	395.067607	18494.670954		20.019552			
14	SSD	btrfs	0.251724	2.3	1.400000	1.4	␣
↪	357.946311	10729.734252		56.053894			
15	SSD	ext2	0.085345	2.0	1.400000	1.4	␣
↪	325.684270	10549.563048		230.214684			
16	SSD	ext3	0.324194	4.8	1.400000	1.4	␣
↪	329.797954	10612.727878		246.229640			
17	SSD	ext4	1.391667	7.3	1.500000	1.5	␣
↪	337.921325	10687.398364		70.066876			
18	SSD	ntfs	2.742155	5.1	3.247685	3.5	␣
↪	317.388026	8631.487349		7132.278231			
19	SSD	vfat	1.273333	2.8	1.400000	1.4	␣
↪	314.693250	3253.701227		28.031059			
20	SSD	xfs	0.227027	2.1	1.400000	1.4	␣
↪	349.935400	10711.527850		72.070390			

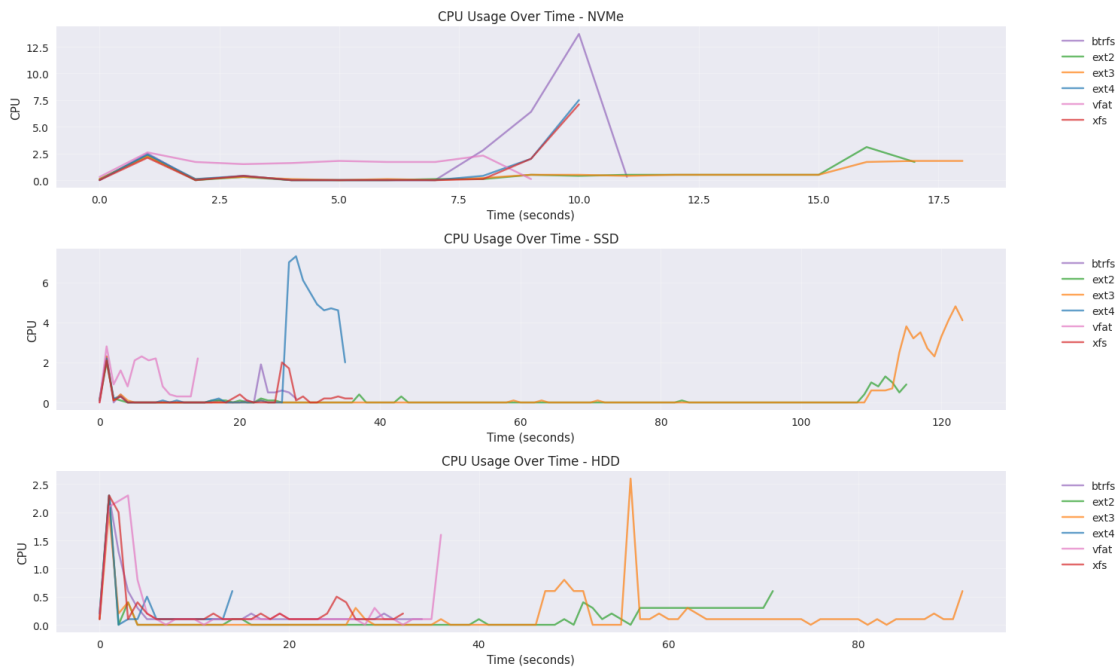
1.7.2 Resource Usage Analysis (excluding NTFS)

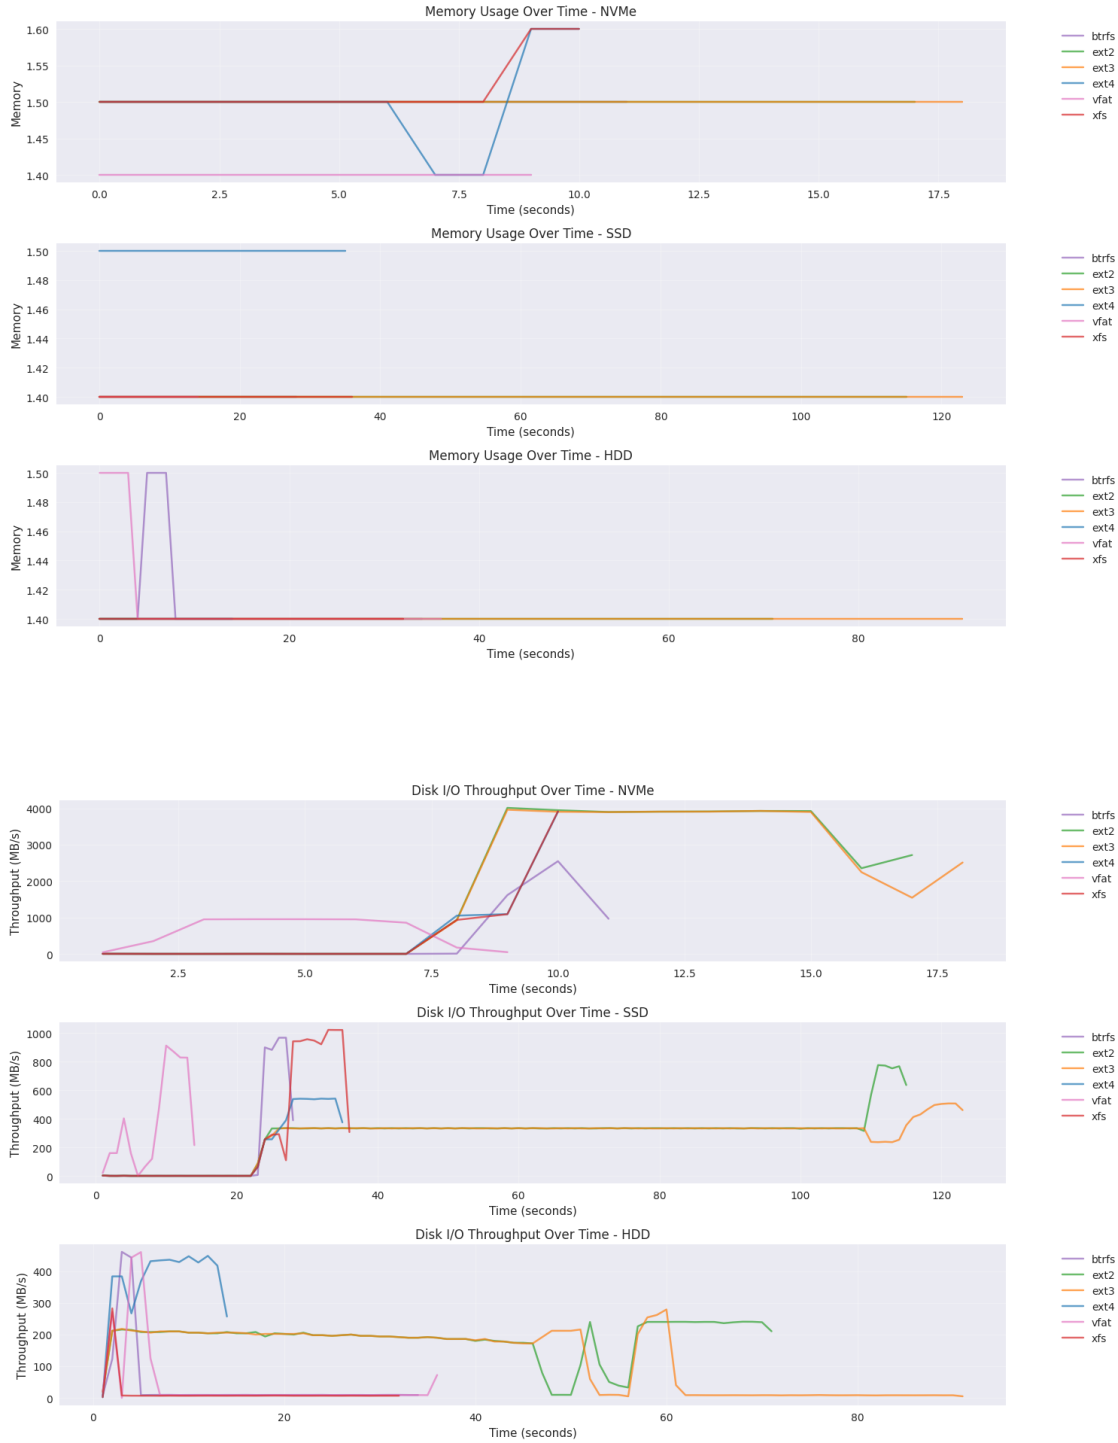
The following analysis excludes NTFS filesystem results due to caching effects that significantly distort the measurements. The significantly longer elapsed time for NTFS tests suggests substantial overhead from formatting and mounting operations before the actual benchmark begins.

```
[9]: # Plot CPU Usage
plot_resource_trends(df[df['filesystem'] != 'ntfs'], 'cpu_percent', 'CPU')

# Plot Memory Usage
plot_resource_trends(df[df['filesystem'] != 'ntfs'], 'memory_percent', 'Memory')

# Plot Disk I/O Throughput
plot_io_throughput(df[df['filesystem'] != 'ntfs'])
```





1.7.3 Resource Usage Matrix Analysis

The resource usage matrix provides a comprehensive visualization of how different filesystems perform across various storage types, showing the temporal patterns of resource utilization during

different test phases.

```
[10]: def plot_resource_matrix(df):
    # Get unique storage types and filesystems
    storage_types = df['storage_type'].unique()
    filesystems = df['filesystem'].unique()

    # Create a figure with subplots matrix
    fig, axes = plt.subplots(len(filesystems), len(storage_types),
                             figsize=(15, 20),
                             squeeze=False)

    # Set title for the entire figure
    fig.suptitle('Resource Usage Patterns by Filesystem and Storage Type',
                 fontsize=16, y=1.02)

    # Test phases in order
    test_phases = ['random_read', 'random_write', 'sequential_read',
    ↪ 'sequential_write']
    phase_colors = ['lightblue', 'lightgreen', 'lightpink', 'lightyellow']

    # Create plots for each combination
    for i, filesystem in enumerate(sorted(filesystems)):
        for j, storage_type in enumerate(storage_types):
            ax = axes[i, j]

            # Get all data for this combination
            mask = (df['storage_type'] == storage_type) & (df['filesystem'] ==
    ↪ filesystem)
            subset = df[mask]

            if len(subset) > 0:
                # Combine monitoring data from all test phases
                all_stats = []
                cumulative_time = 0

                for test_type in test_phases:
                    test_data = subset[subset['test_type'] == test_type]
                    if len(test_data) > 0 and 'monitor_data' in test_data:
    ↪ iloc[0] and 'stats' in test_data.iloc[0]['monitor_data']:
                        stats = pd.DataFrame(test_data.
    ↪ iloc[0]['monitor_data']['stats'])
                        stats['timestamp'] = pd.to_datetime(stats['timestamp'])
                        stats['relative_time'] = range(len(stats))
                        stats['absolute_time'] = stats['relative_time'] +
    ↪ cumulative_time
                        stats['test_phase'] = test_type
```

```

        all_stats.append(stats)
        cumulative_time += len(stats)

    if all_stats:
        combined_stats = pd.concat(all_stats)

        # Calculate I/O throughput
        combined_stats['io_throughput'] =
        ↪(combined_stats['disk_read_bytes'].diff() +
        ↪combined_stats['disk_write_bytes'].diff()) / 1024 / 1024

        # Plot three metrics on the same subplot
        ax2 = ax.twinx()
        ax3 = ax.twinx()
        ax3.spines['right'].set_position(('outward', 60))

        # Plot phase backgrounds
        for phase_idx, phase in enumerate(test_phases):
            phase_data =
            ↪combined_stats[combined_stats['test_phase'] == phase]
            if not phase_data.empty:
                start_time = phase_data['absolute_time'].min()
                end_time = phase_data['absolute_time'].max()
                ax.axvspan(start_time, end_time,
                           color=phase_colors[phase_idx],
                           alpha=0.3,
                           label=phase)

        # Plot metrics
        line1 = ax.plot(combined_stats['absolute_time'],
                        combined_stats['cpu_percent'],
                        color='blue', label='CPU %', alpha=0.6)
        line2 = ax2.plot(combined_stats['absolute_time'],
                        combined_stats['memory_percent'],
                        color='green', label='Memory %', alpha=0.6)
        line3 = ax3.plot(combined_stats['absolute_time'],
                        combined_stats['io_throughput'],
                        color='red', label='I/O (MB/s)', alpha=0.6)

        # Set labels and limits
        ax.set_ylim(0, 100)
        ax2.set_ylim(0, 100)
        ax3.set_ylim(0, combined_stats['io_throughput'].max() * 1.1)

        # Add legend
        lines = line1 + line2 + line3

```



```

        labels = [l.get_label() for l in lines]
        ax.legend(lines + [plt.Rectangle((0,0),1,1, fc=c, alpha=0.
↪3) for c in phase_colors],
                labels + test_phases,
                loc='upper right',
                fontsize='x-small')

    # Set title for each subplot
    ax.set_title(f'{filesystem} - {storage_type}', fontsize=10)

    # Only show x-axis labels for bottom row
    if i == len(filesystems) - 1:
        ax.set_xlabel('Time (s)')

    # Clean up unnecessary ticks
    if i != len(filesystems) - 1:
        ax.set_xticks([])

    # Adjust layout
    plt.tight_layout()
    plt.show()

# Create the matrix plot
plot_resource_matrix(df)

```

Resource Usage Patterns by Filesystem and Storage Type



1.8 7. Comparative Analysis

This section presents a comparative analysis of storage performance metrics across different configurations. It examines three key performance indicators (IOPS, bandwidth, and latency) by comparing both storage types (HDD, SSD, NVMe) and filesystems (ext4, xfs, btrfs, etc.). The analysis includes both tabular data and visual heatmaps to facilitate easy comparison of performance characteristics.

```
[11]: def perform_comparative_analysis(df):  
    """Perform comparative analysis between different configurations."""  
    # Create comparison matrix  
    metrics = ['iops', 'bandwidth', 'latency']  
  
    for metric in metrics:  
        # Compare storage types  
        storage_comparison = pd.pivot_table(  
            df,  
            values=metric,  
            index='test_type',  
            columns='storage_type',  
            aggfunc='mean'  
        )  
  
        print(f"\n{metric.upper()} Comparison by Storage Type:")  
        display(storage_comparison)  
  
        # Compare filesystems  
        fs_comparison = pd.pivot_table(  
            df,  
            values=metric,  
            index='test_type',  
            columns='filesystem',  
            aggfunc='mean'  
        )  
  
        print(f"\n{metric.upper()} Comparison by Filesystem:")  
        display(fs_comparison)  
  
        # Plot comparison heatmaps  
        fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))  
  
        sns.heatmap(storage_comparison, annot=True, cmap='YlOrRd', ax=ax1)  
        ax1.set_title(f'{metric.upper()} by Storage Type')  
  
        sns.heatmap(fs_comparison, annot=True, cmap='YlOrRd', ax=ax2)  
        ax2.set_title(f'{metric.upper()} by Filesystem')  
  
        plt.tight_layout()
```

```
plt.show()


perform_comparative_analysis(df)
```

IOPS Comparison by Storage Type:

storage_type	HDD	NVMe	SSD
test_type			
random_read	5218.303830	100538.353152	21522.412417
random_write	172.426875	30839.360014	12627.320963
sequential_read	401.794746	2118.419987	604.339425
sequential_write	265.055608	1332.734501	458.479528

IOPS Comparison by Filesystem:

filesystem	btrfs	ext2	ext3	ext4	
test_type					
random_read	34150.718689	48464.614005	47448.554914	47416.674581	35829.
random_write	6993.685213	16671.054181	16155.161039	24621.679064	130.
sequential_read	1131.475723	1124.197273	1062.197436	1063.653334	1776.
sequential_write	777.993532	610.454910	569.397336	867.935536	857.




BANDWIDTH Comparison by Storage Type:

storage_type	HDD	NVMe	SSD
test_type			

random_read	20872.714286	4.021530e+05	86089.285714
random_write	689.428571	1.233570e+05	50508.857143
sequential_read	411437.428571	2.169262e+06	618843.285714
sequential_write	271412.142857	1.364703e+06	469475.428571

BANDWIDTH Comparison by Filesystem:

filesystem	btrfs	ext2	ext3	ext4	
↪ ntfs					
↪ vfat					
↪ xfs					
test_type					
random_read	1.366027e+05	1.938580e+05	1.897937e+05	1.896663e+05	1.
↪ 433190e+05	144279.0	1.904163e+05			
random_write	2.797433e+04	6.668367e+04	6.462033e+04	9.848633e+04	5.
↪ 216667e+02	92472.0	5.653733e+04			
sequential_read	1.158631e+06	1.151178e+06	1.087690e+06	1.089181e+06	1.
↪ 819504e+06	0.0	1.159416e+06			
sequential_write	7.966650e+05	6.251050e+05	5.830623e+05	8.887653e+05	8.
↪ 781777e+05	269774.0	8.714947e+05			




LATENCY Comparison by Storage Type:

storage_type	HDD	NVMe	SSD
test_type			
random_read	9.981976e+07	3.793198e+05	1.568308e+06
random_write	2.000650e+08	3.690734e+07	3.698901e+07
sequential_read	1.041555e+08	1.121463e+07	5.007891e+07
sequential_write	1.949310e+08	2.950203e+07	8.141000e+07

LATENCY Comparison by Filesystem:

filesystem	btrfs	ext2	ext3	ext4	
↪ ntfs	vfat	xfs			
test_type					
random_read	3.731838e+07	4.171621e+07	4.207270e+07	3.862764e+07	8.918571e+05
↪	3.947476e+07	3.735570e+07			
random_write	4.959191e+07	9.322743e+07	9.275589e+07	5.205853e+07	2.453643e+08
↪	5.539308e+07	5.085198e+07			
sequential_read	7.134690e+07	7.228924e+07	7.261188e+07	8.070896e+07	1.785679e+07
↪	0.000000e+00	7.123406e+07			
sequential_write	7.580469e+07	1.413831e+08	1.431413e+08	7.700812e+07	3.712815e+07
↪	1.649672e+08	7.420135e+07			



1.9 8. Recommendations and Conclusions

This section presents the final analysis and recommendations based on the benchmark results. The analysis uses a geometric mean approach to calculate overall performance scores, taking into account IOPS, bandwidth, and latency metrics. The recommendations are derived from academic literature on storage benchmarking and statistical analysis of results.

```
[12]: def generate_recommendations(df):
    """Generate recommendations based on benchmark results."""
    # Calculate performance scores using geometric mean of normalized metrics
    # Useful insights from publications:
    # M. Seltzer and K. A. Smith, "Workload-specific file system benchmarks,"
    ↪ 2001. Accessed: May 03, 2025. [Online]. Available: https://www.
    ↪ semanticscholar.org/paper/
    ↪ Workload-specific-file-system-benchmarks-Seltzer-Smith/
    ↪ 686de8280485ad66f0b8037ff61e61b494ce3bc0
```

```

# P. M. Chen and D. A. Patterson, "A new approach to I/O performance
↪evaluation: self-scaling I/O benchmarks, predicted I/O performance,"
↪SIGMETRICS Perform. Eval. Rev., vol. 21, no. 1, pp. 1-12, Jun. 1993, doi: 10.
↪1145/166962.166966.

# P. J. Fleming and J. J. Wallace, "How not to lie with statistics: the
↪correct way to summarize benchmark results," Commun. ACM, vol. 29, no. 3, pp.
↪218-221, Mar. 1986, doi: 10.1145/5666.5673.

# Normalize each metric relative to its mean
normalized_iops = df['iops'] / df['iops'].mean()
normalized_bandwidth = df['bandwidth'] / df['bandwidth'].mean()
normalized_latency = df['latency'].mean() / df['latency'] # Invert latency
↪since lower is better

# Calculate geometric mean as recommended by Fleming & Wallace
# Geometric mean preserves ratios and prevents domination by any single
↪metric
df['performance_score'] = (normalized_iops * normalized_bandwidth *
↪normalized_latency) ** (1/3)

# Group by storage type and filesystem
recommendations = df.groupby(['storage_type', 'filesystem']).agg({
    'performance_score': 'mean',
    'iops': 'mean',
    'bandwidth': 'mean',
    'latency': 'mean'
}).round(2)

print("\nPerformance Scores and Recommendations:")
display(recommendations)

# Plot performance scores
plt.figure(figsize=(10, 6))

sns.barplot(data=df, x='storage_type', y='performance_score',
↪hue='filesystem',
            palette=fs_colors)
plt.title('Overall Performance Score by Storage Type and Filesystem')
plt.xlabel('Storage Type')
plt.ylabel('Performance Score')
plt.legend(title='Filesystem')
plt.tight_layout()
plt.show()

```

```
generate_recommendations(df)
```

Performance Scores and Recommendations:

		performance_score	iops	bandwidth	latency
storage_type	filesystem				
HDD	btrfs	0.09	243.20	117599.75	1.330585e+08
	ext2	0.07	175.52	84171.50	2.143106e+08
	ext3	0.07	175.02	84067.00	2.144025e+08
	ext4	0.09	230.31	110971.75	1.409504e+08
	ntfs	1.36	9388.24	686754.75	7.407539e+07
	vfat	0.04	147.35	30933.00	1.369791e+08
	xfs	0.09	241.13	118222.75	1.344233e+08
NVMe	btrfs	3.19	24226.72	1190281.00	8.778315e+06
	ext2	4.96	41196.32	1184340.00	9.170515e+06
	ext3	4.79	39995.97	1103520.25	9.943538e+06
	ext4	5.63	46855.52	1332587.50	7.463895e+06
	ntfs	1.43	9796.01	731068.50	7.631343e+07
	vfat	5.12	35424.96	261890.75	1.713611e+07
	xfs	4.74	38455.01	1300392.00	7.700020e+06
SSD	btrfs	0.97	7820.49	282023.75	3.370955e+07
	ext2	1.06	8780.89	259106.75	3.798088e+07
	ext3	1.06	8755.49	256287.25	3.859024e+07
	ext4	1.02	8391.63	256014.75	3.788818e+07
	ntfs	1.42	9761.84	713318.75	7.554200e+07
	vfat	1.32	9016.33	87070.00	4.076104e+07
	xfs	1.11	9095.30	289783.25	3.310901e+07

